

Atelier Data Science

Deep learning practice 2
Convolutional Neural Networks


Irina Proskurina

Irina.Proskurina@univ-lyon2.fr

Laboratoire ERIC – Université Lyon 2

Previous Lesson Recap 1


1. Neural Networks -> no feature engineering needed
2. Fully connected layers/Dense/nn.Linear() layers
3. Fully connected neural networks : connect every neuron in one layer to every neuron in the other layer



Previous Lesson Recap 2

MNIST


- Each neuron can detect the presence of a specific set of pixels



Previous Lesson Recap 2

MNIST

- If you shift the digit slightly, the neuron will no longer detect its pattern



Number of parameters


- 784 inputs
- Fully connected layer: 1000 neurons
- Output layer: 10 neurons (one for each class)
- Weights between input and fully connected layers:
 $(784 + 1) * 1000 = 785,000$
- Weights between fully connected and output layers:
 $(1000 + 1) * 10 = 10,010$

Fully connected neural networks for image classification

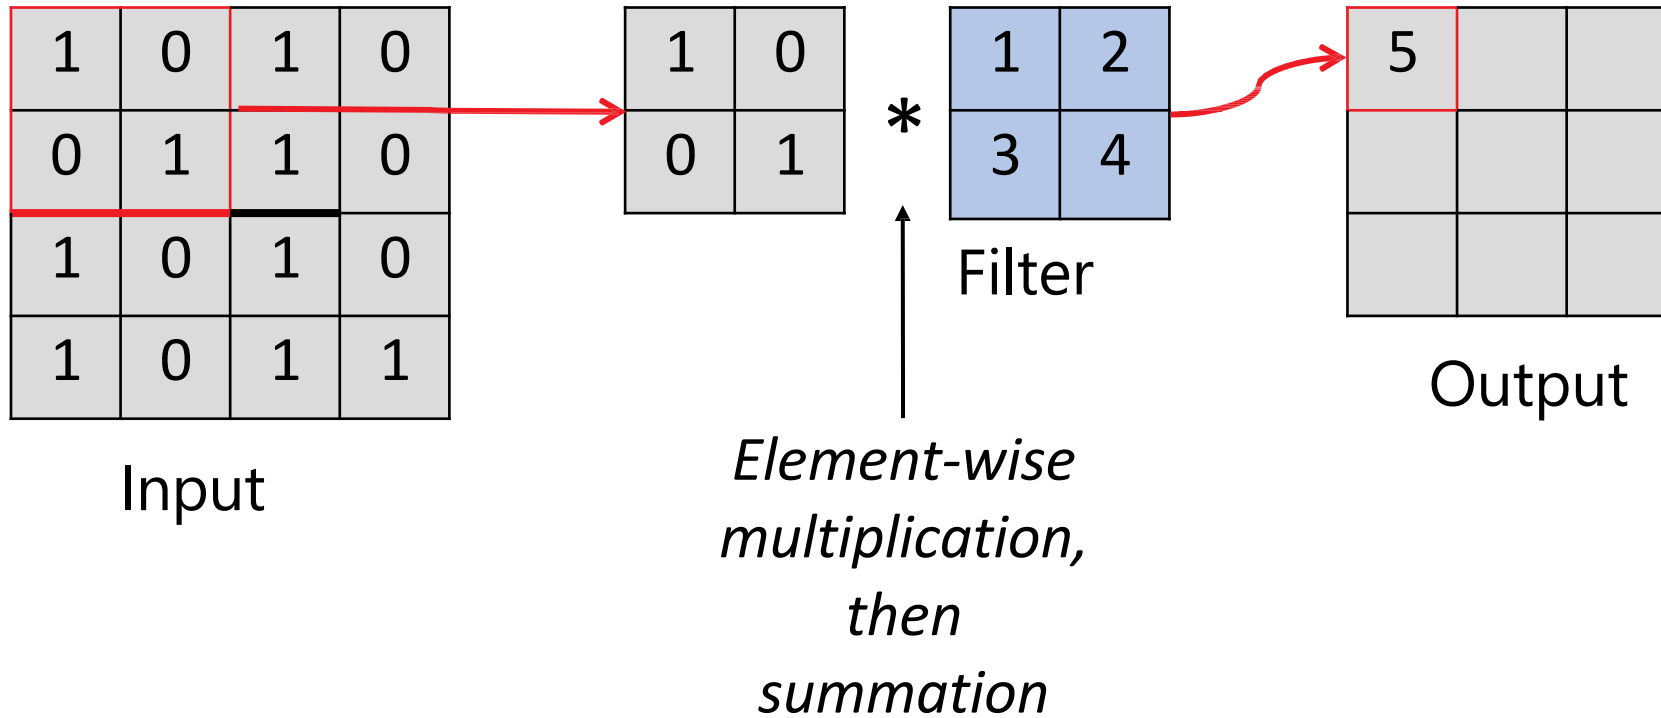
- A lot of parameters
- Prone to overfitting
- Does not consider the specifics of images (shifts, slight changes in shape, etc.)
- One of the best ways to combat overfitting is to reduce the number of parameters

Convolutional neural network

Experiments with the visual cortex



Convolution



Convolution

$$\begin{bmatrix} 1 & 1 \\ 0 & 1 \end{bmatrix} * \begin{bmatrix} 1 & 0 \\ 0 & 1 \end{bmatrix} = \boxed{2}$$

$$\begin{bmatrix} 1 & 1 \\ 1 & 1 \end{bmatrix} * \begin{bmatrix} 1 & 0 \\ 0 & 1 \end{bmatrix} = \boxed{2}$$

$$\begin{bmatrix} 1 & 2 \\ 3 & 0 \end{bmatrix} * \begin{bmatrix} 1 & 0 \\ 0 & 1 \end{bmatrix} = \boxed{1}$$

$$\begin{bmatrix} 0 & 2 \\ 3 & 0 \end{bmatrix} * \begin{bmatrix} 1 & 0 \\ 0 & 1 \end{bmatrix} = \boxed{0}$$

$$\begin{bmatrix} 3 & 0 \\ 0 & 3 \end{bmatrix} * \begin{bmatrix} 1 & 0 \\ 0 & 1 \end{bmatrix} = \boxed{6}$$


$$\begin{bmatrix} 5 & 0 \\ 0 & 5 \end{bmatrix} * \begin{bmatrix} 1 & 0 \\ 0 & 1 \end{bmatrix} = \boxed{10}$$

Convolution



- Detects a pattern in the image, which is defined by a filter
- The stronger the pattern is represented in a particular area of the image, the higher the convolution value will be

Convolution

- The result of convolving an image with a filter is a new image



The convolution maximum is invariant to shifts




Convolutions in computer vision



$$\begin{bmatrix} +1 & 0 & -1 \\ +2 & 0 & -2 \\ +1 & 0 & -1 \end{bmatrix}$$

Horizontal Sobel kernel




Convolutions in computer vision



$$\begin{bmatrix} \bullet 0 & \bullet 0 & \bullet 0 \\ \bullet 0 & \bullet 1 & \bullet 0 \\ \bullet 0 & \bullet 0 & \bullet 0 \end{bmatrix} + \begin{bmatrix} \bullet 0 & \bullet 0 & \bullet 0 \\ \bullet 0 & \bullet 1 & \bullet 0 \\ \bullet 0 & \bullet 0 & \bullet 0 \end{bmatrix} - \frac{1}{9} \begin{bmatrix} \bullet 1 & \bullet 1 & \bullet 1 \\ \bullet 1 & \bullet 1 & \bullet 1 \\ \bullet 1 & \bullet 1 & \bullet 1 \end{bmatrix} = \begin{bmatrix} \bullet 0 & \bullet 0 & \bullet 0 \\ \bullet 0 & \bullet 2 & \bullet 0 \\ \bullet 0 & \bullet 0 & \bullet 0 \end{bmatrix} - \frac{1}{9} \begin{bmatrix} \bullet 1 & \bullet 1 & \bullet 1 \\ \bullet 1 & \bullet 1 & \bullet 1 \\ \bullet 1 & \bullet 1 & \bullet 1 \end{bmatrix}$$

Convolutions in computer vision



$$\frac{1}{9} \begin{bmatrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{bmatrix}$$

Convolution

$$\text{Im}^{out}(x, y) = \sum_{i=-d}^d \sum_{j=-d}^d (K(i, j) \text{Im}^{in}(x + i, y + j) + b)$$


- A pixel in the resulting image depends only on a small region of the input image (local connectivity)
- The weights are the same for all pixels in the output image (shared weights)

Convolution

- Usually, the original image is colored!
- This means that it has multiple channels (R, G, B). Let's consider it in the formula:

$$\text{Im}^{out}(x, y) = \sum_{i=-d}^d \sum_{j=-d}^d \sum_{c=1}^C (K(i, j, c) \text{Im}^{in}(x + i, y + j, c) + b)$$

Convolution



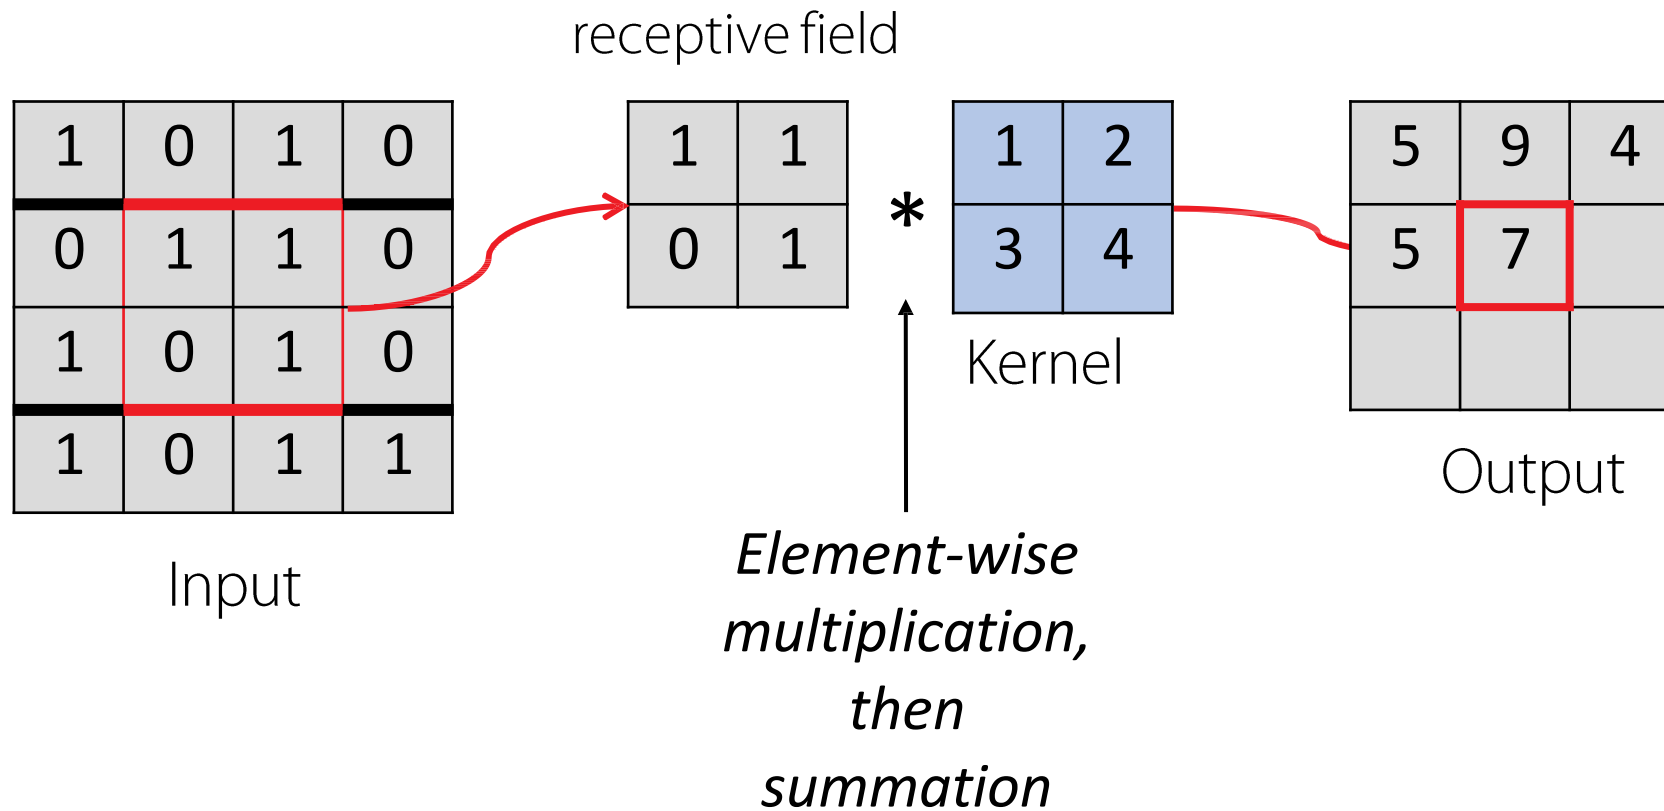
Number of parameters

$$\text{Im}^{out}(x, y, t) = \sum_{i=-d}^d \sum_{j=-d}^d \sum_{c=1}^C (\textcolor{red}{K_t(i, j, c)} \text{Im}^{in}(x + i, y + j, c) + \textcolor{red}{b_t})$$

- Kernel parameters
- $((2d + 1)^2 * C + 1) * T$

Receptive field


Convolution



Receptive field


- Let's take a pixel in the output image
- Which part of the input image does the value in this pixel depend on?

Receptive field



Receptive field: 3 x 3

Receptive field



Receptive field: 5 x 5


Receptive field

Receptive field for 3 x 3 convolution:


- After 1 convolutional layer: 3 x 3
- After 2 convolutional layers: 5 x 5
- After 3 convolutional layers: 7 x 7

We need a lot layers if the image size is 512 x 512


Strides



*




=




$$s = 2$$

Strides




Receptive field: 7 x 7

Dilated convolutions




$l = 2$


Dilated convolutions



$l = 1$





$l = 2$



$l = 4$

The convolution maximum is invariant to shifts



Pooling

1	0	2	1	0	0
0	1	3	2	1	2



1	3	2

Max-pooling with kernel 2x2

Pooling

- Splits the image into $n \times m$ sections applying some function (usually a maximum)
- Significantly reduces the size of the image (which means it increases the field of perception of the following layers)
- Has no parameters


Why we need to know all this?

- It is important to ensure that the last convolutional layers have a perceptual field size comparable to the entire image

Convolution

- If you apply convolution, the output image will be smaller than the input


Convolutions



Valid mode

- When counting convolutions, the pixels at the edges do not have a big impact on the result

We will not see that the filter has a good response when placing the center at this pixel



Zero padding

0	0	0	0	0	0	0	0
0							0
0							0
0							0
0							0
0							0
0							0
0	0	0	0	0	0	0	0

[illegible]

Zero padding

- We add zeros along the boundaries so that the convolution calculated after this in valid mode gives an image of the same size as the original one
- There is a risk that the model will learn to understand where the edges are in the image - we may lose invariance

Reflection padding

[illegible]

[illegible]

Reflection padding

- Can't easily find image edges
- But now the model can begin to find specular reflections and select filters for them

Replication padding

[illegible]

[illegible]

Replication padding

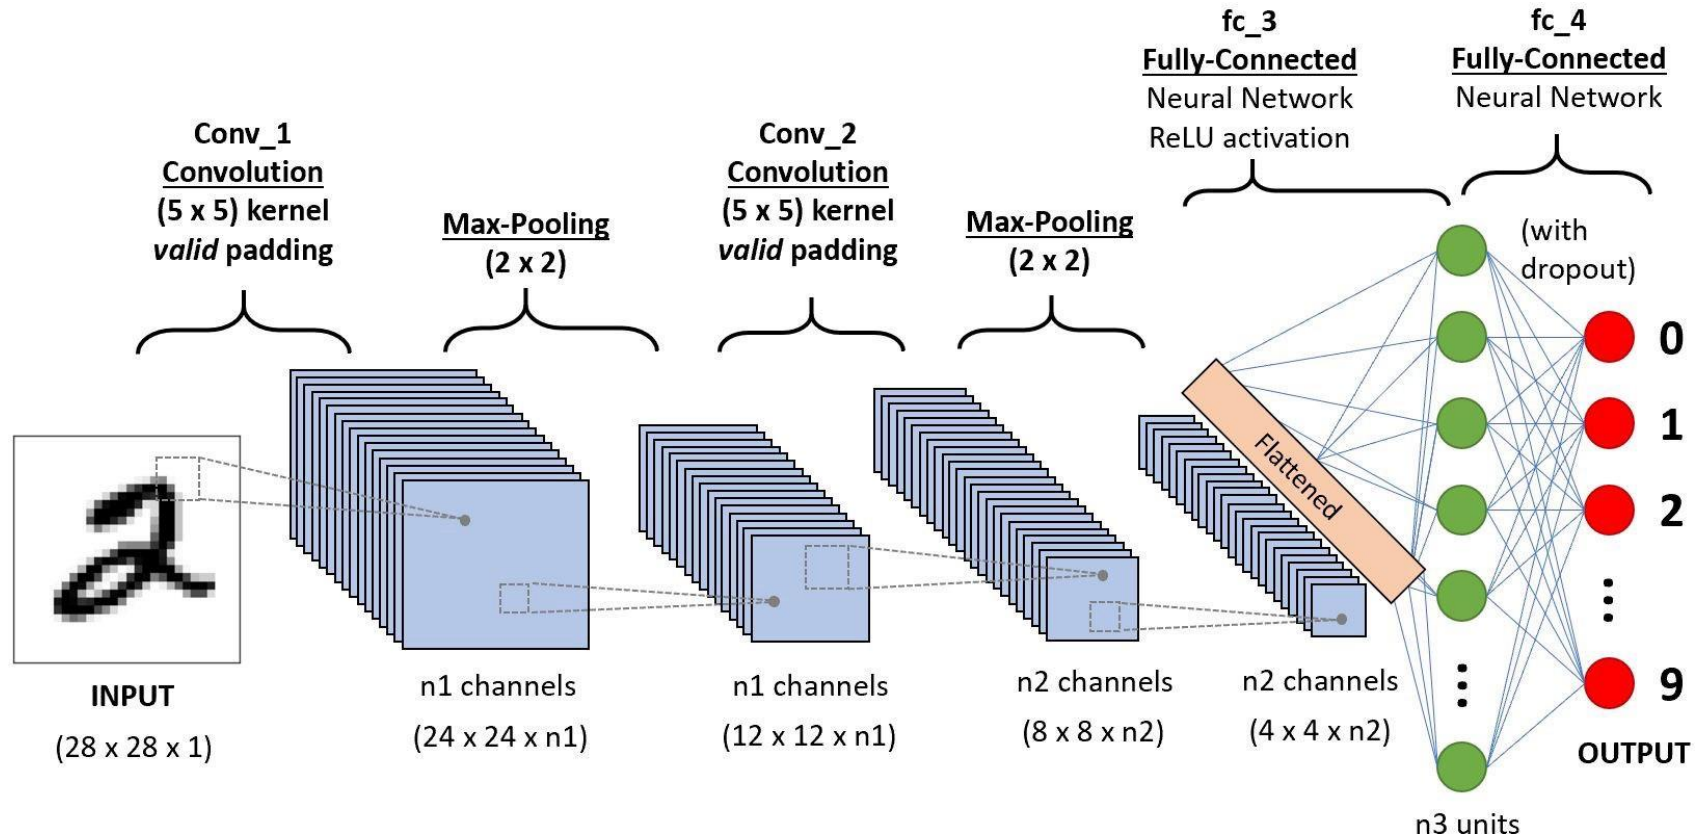
- The pixel on the border is equal to the nearest pixel from the image
- The model can still adjust to the patterns that arise from such padding

Summary

- Padding allows you to control the size of the output images
- Padding allows you to take into account objects on the edges
- Different types of padding allow different methods of retraining for edges

Convolutional Neural Networks


Architecture




Architecture

- Convolution->linear layer>pooling or convolution->non-linear layer
- flattening of the output
- Fully-connected layer

LeNet



AlexNet

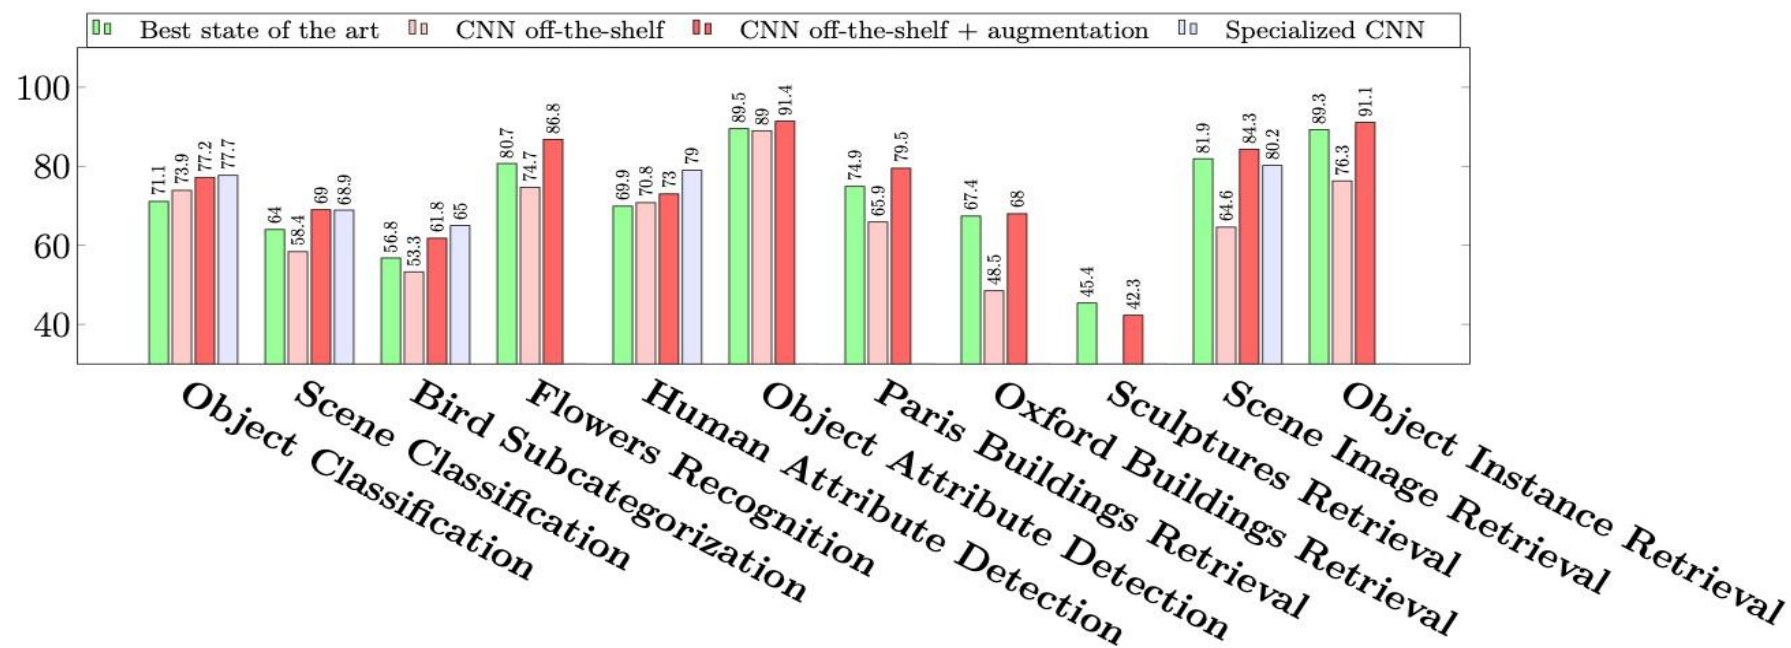
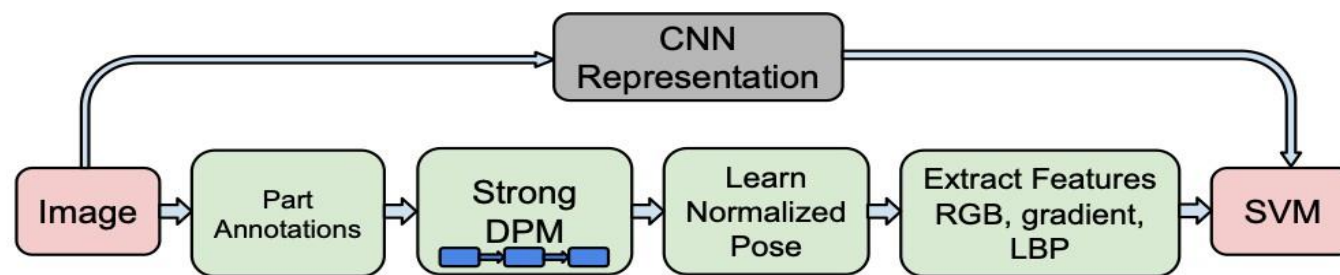


<http://papers.nips.cc/paper/4824-imagenet-classification-with-deep-convolutional-neural-networks.pdf>

Image representation (embedding) from the last layers

- Important observation: the outputs of fully connected layers serve as good feature representations of images and are valuable in many tasks
- For instance, they can be utilized in tasks like searching for similar images

Last layer embeddings





Last layer embeddings




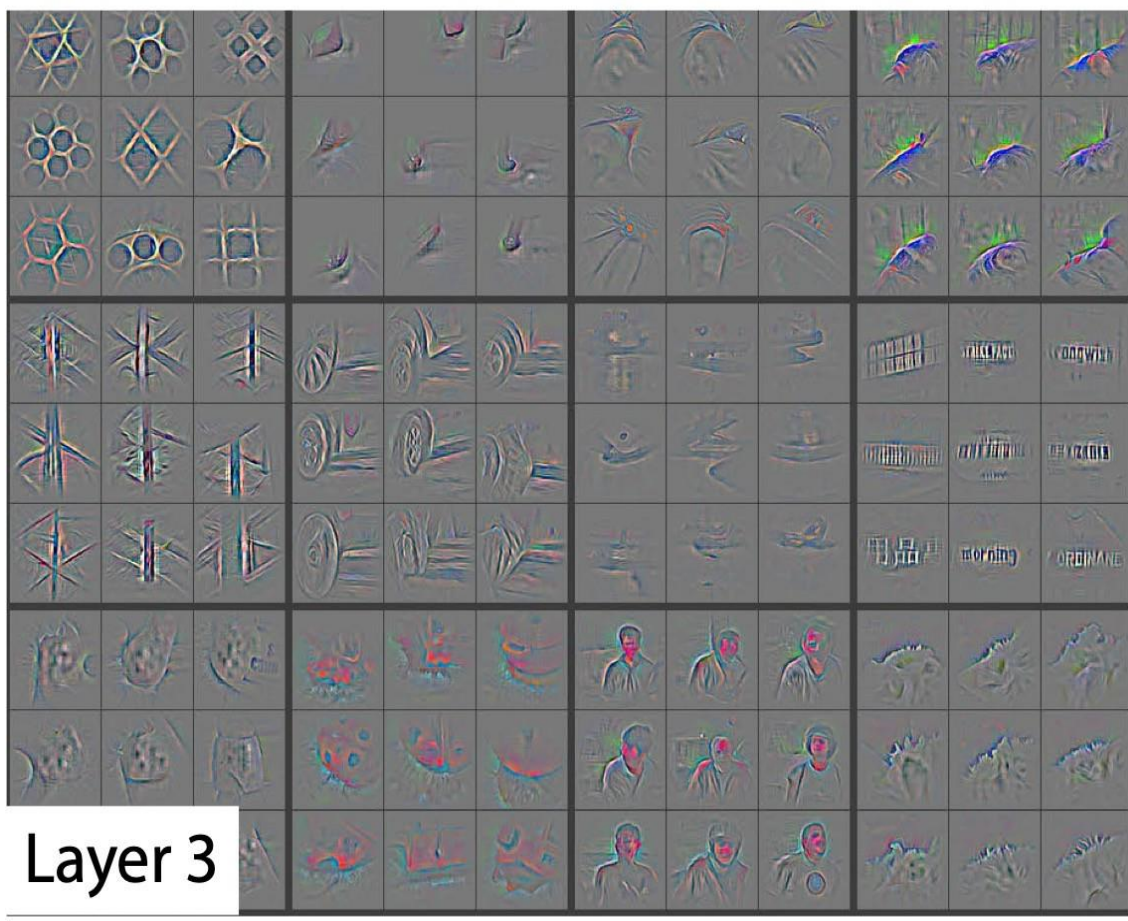
Layer 1



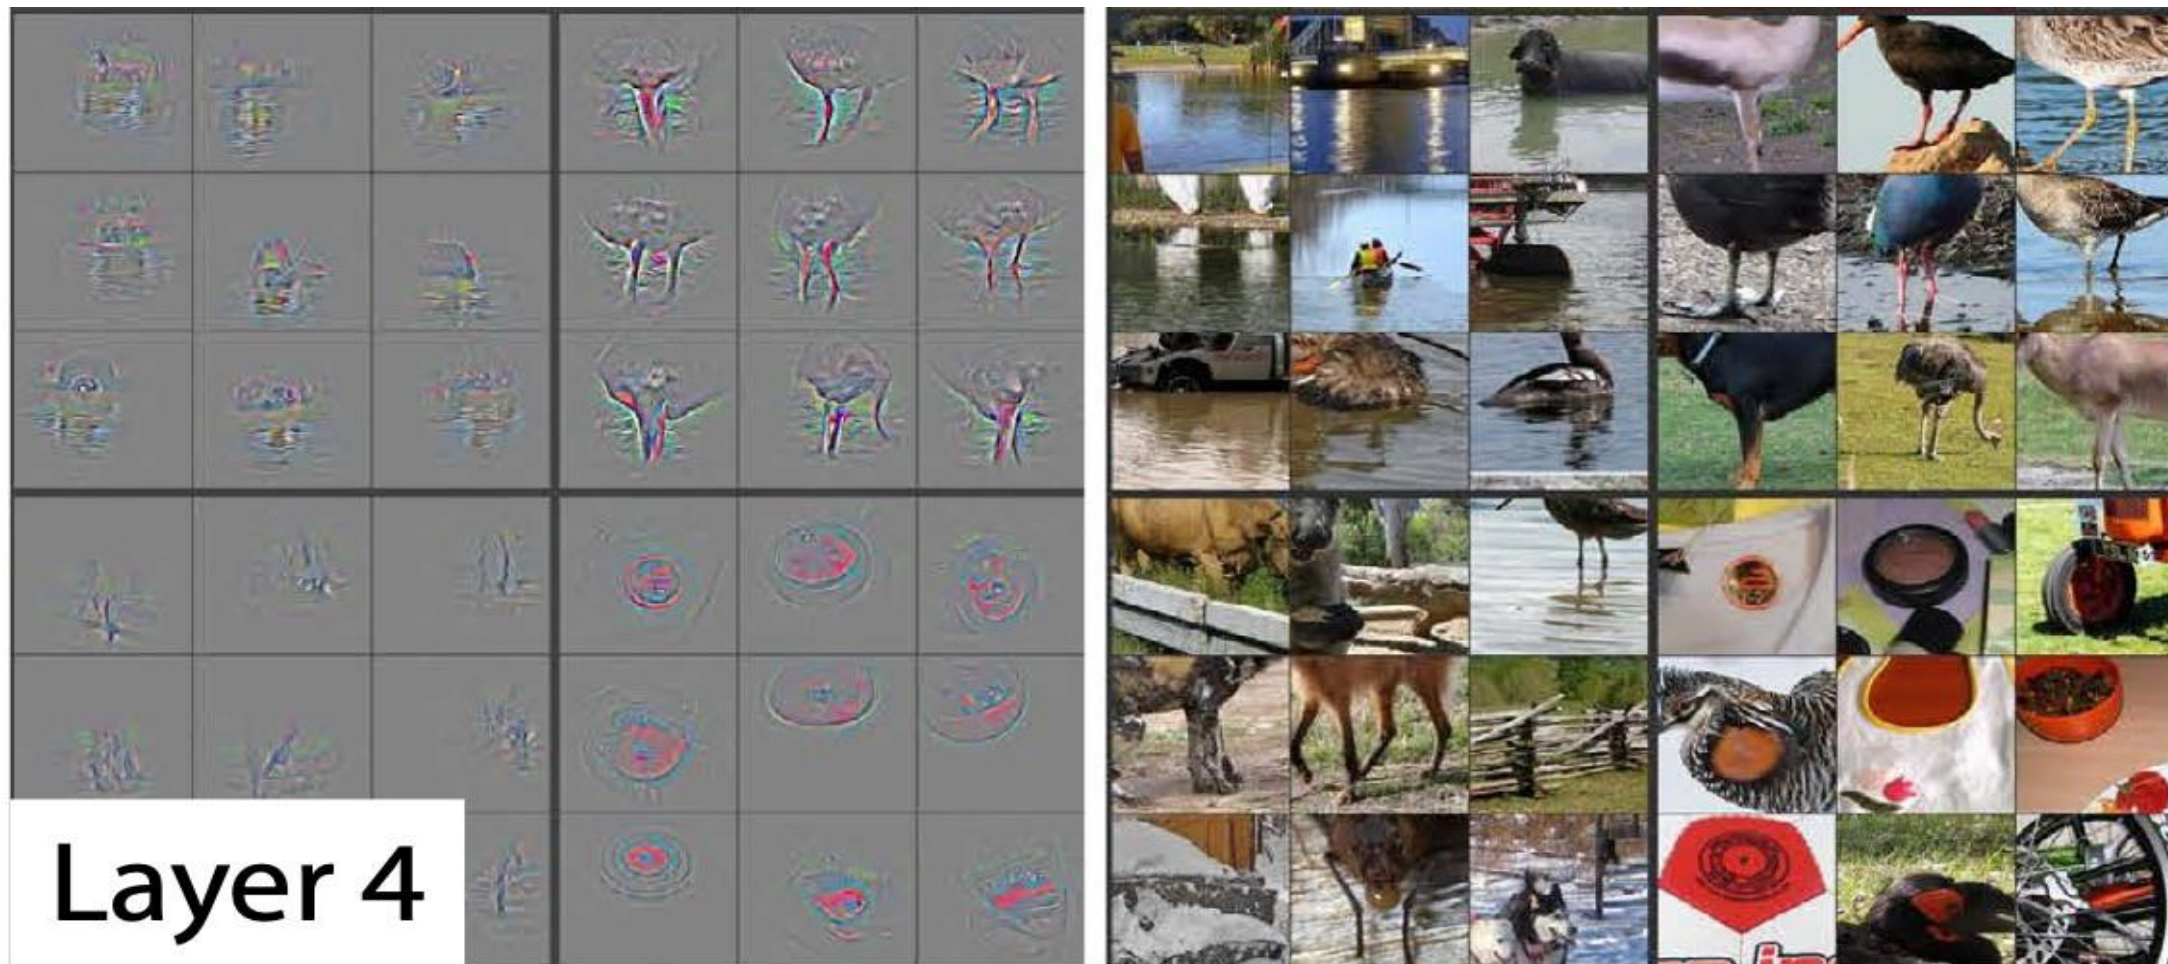
Last layer embeddings



Last layer embeddings



Last layer embeddings



Last layer embeddings

